

# TDT4225 - Assignment 3

*Group 26*

Ole Jonas Liahagen  
Martin Johannes Nilsen

22nd October 2021

---

# Introduction

In this assignment we were to insert and process data from the Geolife GPS Trajectory dataset. The data given to us consisted of 182 unique users, each with their own registered activities. These activities each consists of a list of tracking points containing GPS coordinates and timestamps. Activities could either be labeled with a transportation mode or not at all. To solve the assignment, we have used a combination of database queries and Python to get the data desired for answering each task.

Only minimal changes to the Python script were required for inserting data to a mongodb database instead of the MySQL database in the former assignment. We had already created dataclasses, and as classes in Python are stored as dictionaries, we could easily insert these to the db with a few adjustments of the code. After we had built the database structure and inserted data, we split up and solved the assignments separately, with the possibility to ask each other questions and share some thoughts to check that we agreed on the solutions. We used GitHub<sup>1</sup> as a version control system for our code, and used Overleaf to collaboratively write and compile this L<sup>A</sup>T<sub>E</sub>X report.

---

<sup>1</sup><https://github.com/MartinJohannesNilsen/tdt4225-assignment3.git>

---

# Results

## Part 1

The top 10 documents in each collection:

```
User:
{'_id': 000, 'has_labels': False, 'activities': [ ... ]}
{'_id': 001, 'has_labels': False, 'activities': [ ... ]}
{'_id': 002, 'has_labels': False, 'activities': [ ... ]}
{'_id': 003, 'has_labels': False, 'activities': [ ... ]}
{'_id': 004, 'has_labels': False, 'activities': [ ... ]}
{'_id': 005, 'has_labels': False, 'activities': [ ... ]}
{'_id': 006, 'has_labels': False, 'activities': [ ... ]}
{'_id': 007, 'has_labels': False, 'activities': [ ... ]}
{'_id': 008, 'has_labels': False, 'activities': [ ... ]}
{'_id': 009, 'has_labels': False, 'activities': [ ... ]}
```

The first 10 documents in the user collection. Activities is a list of all activity ids.

```
Activity:
{'_id': 0, 'user_id': '000', 'transportation_mode': None, 'start_date_time': datetime.datetime(2008, 10, 23, 2, 53, 4), 'end_date_time': datetime.datetime(2008, 10, 23, 11, 11, 12)}
{'_id': 1, 'user_id': '000', 'transportation_mode': None, 'start_date_time': datetime.datetime(2008, 10, 24, 2, 9, 59), 'end_date_time': datetime.datetime(2008, 10, 24, 2, 47, 6)}
{'_id': 2, 'user_id': '000', 'transportation_mode': None, 'start_date_time': datetime.datetime(2008, 10, 26, 15, 44, 7), 'end_date_time': datetime.datetime(2008, 10, 26, 15, 4, 7)}
{'_id': 3, 'user_id': '000', 'transportation_mode': None, 'start_date_time': datetime.datetime(2008, 10, 27, 11, 54, 49), 'end_date_time': datetime.datetime(2008, 10, 27, 12, 5, 54)}
{'_id': 4, 'user_id': '000', 'transportation_mode': None, 'start_date_time': datetime.datetime(2008, 10, 28, 0, 38, 26), 'end_date_time': datetime.datetime(2008, 10, 28, 5, 3, 42)}
{'_id': 5, 'user_id': '000', 'transportation_mode': None, 'start_date_time': datetime.datetime(2008, 10, 29, 9, 21, 38), 'end_date_time': datetime.datetime(2008, 10, 29, 9, 30, 28)}
{'_id': 6, 'user_id': '000', 'transportation_mode': None, 'start_date_time': datetime.datetime(2008, 10, 29, 9, 30, 38), 'end_date_time': datetime.datetime(2008, 10, 29, 9, 46, 43)}
{'_id': 7, 'user_id': '000', 'transportation_mode': None, 'start_date_time': datetime.datetime(2008, 11, 3, 10, 13, 36), 'end_date_time': datetime.datetime(2008, 11, 3, 10, 16, 1)}
{'_id': 8, 'user_id': '000', 'transportation_mode': None, 'start_date_time': datetime.datetime(2008, 11, 3, 23, 21, 53), 'end_date_time': datetime.datetime(2008, 11, 4, 3, 31, 8)}
{'_id': 9, 'user_id': '000', 'transportation_mode': None, 'start_date_time': datetime.datetime(2008, 11, 10, 1, 36, 37), 'end_date_time': datetime.datetime(2008, 11, 10, 3, 46, 12)}
```

The first 10 documents in the activity collection

```
TrackPoint:
{'_id': 0, 'latitude': '39.984702', 'longitude': '116.318417', 'altitude': '492', 'date_days': '39744.1201851852', 'date_time': datetime.datetime(2008, 10, 23, 2, 53, 4), 'activity_id': 0}
{'_id': 1, 'latitude': '39.984683', 'longitude': '116.31845', 'altitude': '492', 'date_days': '39744.1202546296', 'date_time': datetime.datetime(2008, 10, 23, 2, 53, 10), 'activity_id': 0}
{'_id': 2, 'latitude': '39.984686', 'longitude': '116.318417', 'altitude': '492', 'date_days': '39744.12031325', 'date_time': datetime.datetime(2008, 10, 23, 2, 53, 15), 'activity_id': 0}
{'_id': 3, 'latitude': '39.984688', 'longitude': '116.318385', 'altitude': '492', 'date_days': '39744.1203783704', 'date_time': datetime.datetime(2008, 10, 23, 2, 53, 20), 'activity_id': 0}
{'_id': 4, 'latitude': '39.984655', 'longitude': '116.318263', 'altitude': '492', 'date_days': '39744.1204282407', 'date_time': datetime.datetime(2008, 10, 23, 2, 53, 25), 'activity_id': 0}
{'_id': 5, 'latitude': '39.984611', 'longitude': '116.318026', 'altitude': '493', 'date_days': '39744.1204861111', 'date_time': datetime.datetime(2008, 10, 23, 2, 53, 30), 'activity_id': 0}
{'_id': 6, 'latitude': '39.984608', 'longitude': '116.317761', 'altitude': '493', 'date_days': '39744.1205439815', 'date_time': datetime.datetime(2008, 10, 23, 2, 53, 35), 'activity_id': 0}
{'_id': 7, 'latitude': '39.984563', 'longitude': '116.317517', 'altitude': '496', 'date_days': '39744.1206018519', 'date_time': datetime.datetime(2008, 10, 23, 2, 53, 40), 'activity_id': 0}
{'_id': 8, 'latitude': '39.984539', 'longitude': '116.317294', 'altitude': '500', 'date_days': '39744.1206597222', 'date_time': datetime.datetime(2008, 10, 23, 2, 53, 45), 'activity_id': 0}
{'_id': 9, 'latitude': '39.984606', 'longitude': '116.317065', 'altitude': '505', 'date_days': '39744.1207175926', 'date_time': datetime.datetime(2008, 10, 23, 2, 53, 50), 'activity_id': 0}
```

The first 10 documents in the trackpoint collection

---

## Part 2

### Task 1

The amount of users, activities and trackpoints in the dataset after inserted to the database:

Task 1		
n_users	n_activities	n_trackpoints
-----	-----	-----
182	16048	9681756

### Task 2

The *average*, *minimum* and *maximum* number of activities per user. It is important to notice that some users have tracked activities, but because we exclude all activities with more than 2500 trackpoints, some users end up having zero activities.

Task 2		
average	maximum	minimum
-----	-----	-----
88.1758	2102	0

### Task 3

Top 10 users with the highest number of activities:

Task 3	
Id	Count
----	-----
128	2102
153	1793
025	715
163	704
062	691
144	563
041	399
085	364
004	346
140	345

---

## Task 4

The number of users that have started the activity in one day and ended the activity the next day:

```
Task 4
Users with multi-day activities
-----
98
```

## Task 5

Activities that are registered multiple times, i.e. the same start- and end-time, but different ids. The dataset includes zero duplicated activities.

```
Task 5
activities registred multiple times
-----
```

## Task 6

Users who have been close to the infected person at position (39.97548, 116.33031) the 24th of august 2008, at 15:38.

```
Task 6
Users in contact: ['073']
```

## Task 7

In this task we were requested to find all users that have never taken a taxi. The table includes a column named *labeled* to show if the concerning user is labelled or not. If the user is not labelled, they will obviously never have a labelled taxi ride registered. The result set contains 172 rows, which indicates that only 10 users have taken a taxi and labelled the ride (182-172). Out of the 172 who have never taken a taxi, 69 users have labelled their data.

---

Task 7		
Users that have never taken a taxi		labeled
-----		-----
	020	True
	021	True
	052	True
	053	True
	056	True
	059	True
	060	True
	064	True
	065	True
	067	True
	068	True
	069	True
	073	True
	075	True
	076	True
	081	True
	082	True
	084	True
	086	True
	087	True
	088	True
	089	True
	091	True
	092	True
	096	True
	097	True
	100	True
	101	True
	102	True
	104	True
	105	True
	106	True
	107	True
	108	True
	110	True
	112	True
	114	True
	115	True
	116	True
	117	True
	118	True
	124	True
	125	True
	126	True
	129	True
	136	True
	138	True
	139	True
	141	True
	144	True
	147	True
	153	True
	154	True
	161	True
	167	True
	170	True
	174	True
	175	True
	179	True

---

```
000 False
001 False
002 False
003 False
004 False
005 False
006 False
007 False
008 False
009 False
011 False
012 False
013 False
014 False
015 False
016 False
017 False
018 False
019 False
022 False
023 False
024 False
025 False
026 False
027 False
028 False
029 False
030 False
031 False
032 False
033 False
034 False
035 False
036 False
037 False
038 False
039 False
040 False
041 False
042 False
043 False
044 False
045 False
046 False
047 False
048 False
049 False
050 False
051 False
054 False
055 False
057 False
061 False
063 False
066 False
070 False
071 False
072 False
074 False
077 False
079 False
083 False
090 False
093 False
094 False
095 False
```

---

099	False
103	False
109	False
113	False
119	False
120	False
121	False
122	False
123	False
127	False
130	False
131	False
132	False
133	False
134	False
135	False
137	False
140	False
142	False
143	False
145	False
146	False
148	False
149	False
150	False
151	False
152	False
155	False
156	False
157	False
158	False
159	False
160	False
162	False
164	False
165	False
166	False
168	False
169	False
171	False
172	False
173	False
176	False
177	False
178	False
180	False
181	False



---

## Task 8

All types of transportation modes and how many distinct users that have used them:

Task 8	
transportation_mode	Count
airplane	1
bike	19
boat	1
bus	12
car	8
run	1
subway	4
taxi	10
train	2
walk	31

## Task 9

a) The year and month with the most activities:

Task 9a	
Year with most activities	Month with most activities in this year
2008	11

b) The two users with the most activities in the month and year found from 9a (11-2008), and how many recorded hours they have. The user with the second most activities has more registered hours than the user with the most registered activities.

Task 9b		
uid	n_activities	hours
062	130	47.3136
128	75	68.2211

---

## Task 10

The total distance (in km) walked in 2008, by user with id=112. This distance seems small, but we think it is correct after going over our code and queries. Looking over the trackpoints associated with user 112, it seems they have not moved much around by foot this year. This might be due to our method of labeling the data, so that certain walks were not accepted by our filtering. However, this is only speculation.

```
Task 10
User_id: 112
Distance in 2008: 1.3497848643311852
```

## Task 11

The top 20 users who have gained the most altitude meters:

```
Task 11
uid      m gained
-----
128      650951.9973
153      554960.6231
4         332036.3184
41        240768.8657
3         233663.6424
85        217643.3849
163       205274.3705
62        181693.2917
144       179441.5245
30        175679.7096
39        146703.5928
84        131161.2312
0         121504.8624
2         115198.2456
167       112974.1546
25        109158.5726
37         99234.5894
140       94846.2994
126       83025.8358
17        62581.3531
```

---

## Task 12

In the figures below you will find the users who have invalid activities, and the number of invalid activities per user. An invalid activity is when the time difference between two trackpoints is 5 minutes or more.

Task 12	
uid	Invalid activities
----	-----
0	101
1	45
2	98
3	179
4	219
5	45
6	17
7	30
8	16
9	31
10	50
11	32
12	43
13	29
14	118
15	46
16	20
17	129
18	27
19	31
20	20
21	7
22	55
23	11
24	27
25	263
26	18
27	2
28	36
29	25
30	112
31	3
32	12
33	2
34	88

---

35	23
36	34
37	100
38	58
39	147
40	17
41	201
42	55
43	21
44	32
45	7
46	13
47	6
48	1
49	0
50	8
51	36
52	44
53	7
54	2
55	15
56	7
57	16
58	13
59	5
60	1
61	12
62	249
63	8
64	7
65	26
66	6
67	33
68	139
69	6
70	5
71	29
72	2
73	18
74	19
75	6
76	8
77	3
78	19
79	2
80	6
81	16
82	27
83	15
84	99
85	184
86	5
87	3
88	11
89	40
90	3
91	63
92	101
93	4
94	16
95	4
96	35
97	14
98	5
99	11
100	3
101	46
102	13
103	24
104	97
105	9
106	3
107	1
108	5

---

109	3
110	17
111	26
112	67
113	1
114	3
115	58
116	0
117	3
118	3
119	22
120	0
121	4
122	6
123	3
124	4
125	25
126	105
127	4
128	720
129	6
130	8
131	10
132	3
133	4
134	31
135	5
136	6
137	0
138	10
139	12
140	86
141	1
142	52
143	0
144	157
145	5
146	7
147	30
148	0
149	0
150	16
151	1
152	2
153	557
154	14
155	30
156	0
157	9
158	9
159	5
160	0
161	7
162	9
163	233
164	6
165	2
166	2
167	134
168	19
169	9
170	2
171	3
172	9
173	5
174	54
175	4
176	8
177	0
178	0
179	28
180	2
181	14

---

## Discussion

To start off the assignment, we discussed how to structure our data to make our queries the most effective. The decision landed on every user having a list of activity ids as a one to many relation. Each trackpoint also has a reference to the activity that it belongs to, but not the other way around. The activities do not have a reference to each of their trackpoints. One could argue that including such a list could be beneficial due to our limit of 2500 trackpoints per activity, however, we did not see the need to add such a relation. It would just increase complexity and inserts needed. This leads to our database model being quite similar to the one for the MySQL assignment, with the exception of the arrays of activity ids in the user documents. Our choice of modeling the database this way may be influenced by our previous experience with relational databases and little experience with document databases. There are possibilities of denormalisation here, such as adding data from the trackpoints to the activities. There might even be a more compelling case for doing this rather than putting activity ids in the user documents. If one were to look at the user documents as a complete user object, one would not need activities to define a user, however, it could be argued that an activity needs a list of coordinates to track where the activity found place. There really is no right or wrong here, and it is mostly up to preference. Our preference leaned towards the approach we have implemented, and that is how our design came to be.

The design we chose made the tasks quite similar to the ones from the MySQL assignment, with some differences here and there. Seeing as joins are considered slow in document databases, these were mostly avoided in favor of doing application level joins on retrieved documents instead.

It is also worth mentioning that a teaching assistant fortunate enough to read both of our group assignments with a very good memory would notice that our answers to tasks 11 and 12 differ from the ones given in assignment 2. This is due to us finding and correcting some errors that were present in our previous delivery. Our old answer for task 11 did not perform a conversion from feet to meters, resulting in a very large amount of altitude meters gained for several users. This is now resolved. Task 12 was missing a way to skip past an activity's remaining trackpoints if said activity already had been classified as invalid. This resulted in an invalid activity being able to register as several invalid activities instead of just one. This check is now added, and the results should be correct now.

All in all, working with a document database was less of a hassle than expected, albeit a little tricky to begin with. But after some trial and error, we feel like we got the knack of it, and we are pretty certain we deliver a new and more correct solution this time around!

---

## MySQL vs MongoDB

MongoDB, having a document-collection database format, differ from MySQL in multiple ways. The first obvious difference is the use of collections and documents, instead of tables and rows. This is more reminiscent of how objects are stored in code and can benefit comprehension of the data stored in the database. This also makes it easy to insert data from code to the database, since the JSON format is the standard way to store data in MongoDB. In python we used dictionaries to send data to the database, which is then easily translated to a document in the JSON format. Doing this in MySQL without middleware is a bit more cumbersome, requiring the user to insert data into all the concerning tables. If there is a large object to be inserted, this could span many tables if the database is normalised. The same problem arises when retrieving objects from a relational model. A user would have to assemble the object from each table before returning it for use in code.

The lack of an enforced schema is another thing differentiating MongoDB from MySQL. The flexibility may be a significant advantage for some scenarios, meanwhile in others the guarantee of a given format may be wanted or even required. It is worth mentioning, that while not mandatory, MongoDB does have the opportunity to enforce schemas with special flags in insertions and creations of collections.

After working with both types of systems, we do not feel there is a "better" or "worse" approach to the database model overall. Each system has it's up- and downsides. MongoDB is great for easy access to more complex objects, but slower for lookups in lots of tables. This could however, be overcome by denormalising data. Unfortunately this might lead to more complexity in making the database itself and maintaining it. MySQL keeps this simple and with a properly normalised database, there should be no reason to update fields in any other table than the one you are using at any given moment. Joins in a relational database are also often much faster due to built-in and optimised support for such operations.

As stated earlier - the two models are two different approaches at database systems, that work well in their own respective use-cases. For this assignment, we did not feel there was a "better" version among the two.